

Automatic Image Analysis Exercise 3

# **The Generalised Hough Transformation**


(Due 23.06.2017)

# Mid-term Exam


- Friday, **09.06.2017, 10:15am, room H 3010**
- In place of an exercise
- Duration: ca. 30 min
- No grade, but pass is necessary to take part at the final exam
- Topics from lecture **and** exercise
- Questions in English, answers in English or German
- No books, no calculator, no script, no paper, ...

# $(\rho, \theta)$ Invariance

A line in the image plane is defined by parameters  $(\rho, \theta)$ :



In parameter space over  $(\rho, \theta)$ , the line collapses to a point:



Invariant for all pixels on the line:

$$\begin{pmatrix} x \\ y \end{pmatrix} \cdot \begin{pmatrix} \cos(\theta) \\ \sin(\theta) \end{pmatrix} = \rho$$

## Line Detection

- 0) Compute  $(\rho, \theta)$  at each line-pixel
- 1) Increment the parameter space at coordinate  $(\rho, \theta)$
- 2) A line leads to a pronounced maximum



# Radon Transformation

**Problem:**  $(\rho, \theta)$  are unknowns



- Infinite number of potential lines intersect at every point

**Solution:**


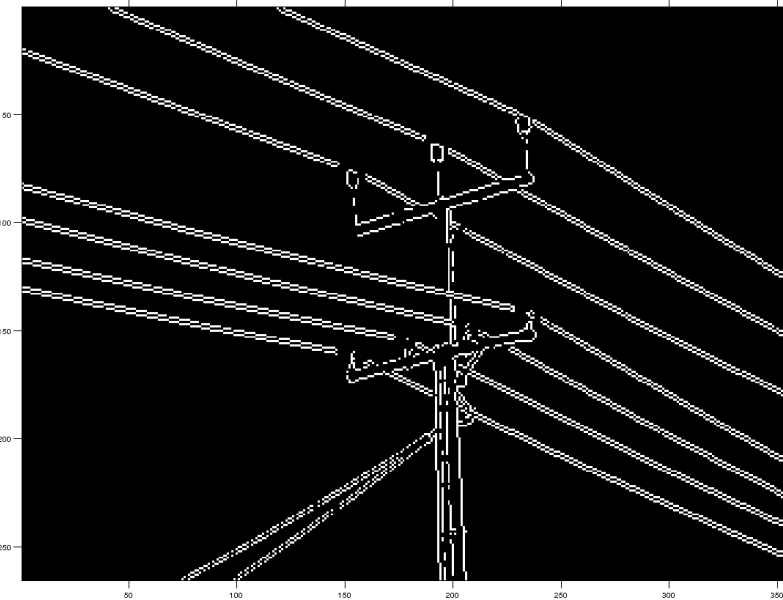
- Each line through a point  $(x, y)$  is defined by fixing its direction  $\theta$ .
- $\rho$  can be computed from  $x, y$  and  $\theta$ .
- For all directions  $\theta$ : Compute  $\rho$  and increase element  $(\rho, \theta)$  in parameter space




# Radon Transformation



# Radon Transformation




# Radon Transformation



# Hough Line Detector

- Alternative (more intelligent) solution: Use gradient information
- The local gradient direction fixes  $\theta$ , and  $(\rho, \theta)$  can be calculated directly
- It is no longer necessary to consider every possible line through each pixel!





# Hough Line Detector

- Alternative (more intelligent) solution: Use gradient information
- The local gradient direction fixes  $\theta$ , and  $(\rho, \theta)$  can be calculated directly
- It is no longer necessary to consider every possible line through each pixel!






Image  
 $B(x, y)$





Gradient in x direction  
 $G_x(x, y) = \frac{\partial}{\partial x} B(x, y)$



Gradient in y direction  
 $G_y(x, y) = \frac{\partial}{\partial y} B(x, y)$

# Hough Line Detector

- Each pixel  $(x,y)$  is associated with a gradient vector  $(G_x(x,y), G_y(x,y))$ .




- The angle  $\theta_G(x,y)$  is easy to obtain, e.g. in C++:

$$\theta_G(x,y) = \text{atan2}(G_y(x,y), G_x(x,y))$$

- Once  $\theta_G(x,y)$  is known, coordinates  $(\rho, \theta)$  in parameter space are uniquely determined:

$$\begin{pmatrix} x \\ y \end{pmatrix} \cdot \begin{pmatrix} \cos(\theta) \\ \sin(\theta) \end{pmatrix} = \rho$$

# Hough Linien-detektor




- 1) Compute gradients in  $x$  and  $y$
- 2) Determine  $\rho$  und  $\theta$  using the local gradient direction
- 3) Increase  $H_L(\rho, \theta)$  by the local gradient magnitude:

$$H_L(\rho, \theta) \leftarrow H_L(\rho, \theta) + \sqrt{G_x(x, y)^2 + G_y(x, y)^2}$$

A line in  $I(x, y)$  produces a pronounced maximum  $H_L(\rho, \theta)$ .

# Hough Liniendetektor



# Generalised Hough Transformation


- An image  $I(x,y)$  is processed to discover instances of an object
- The object of interest is known from an example image  $O(x,y)$
- Processing is once again concerned with accumulating „votes“ in a parameter space  $H_G(x,y,\theta,s)$ .



Image  $I(x,y)$



Object  $O(x,y)$





$H_G(x,y,\theta,s)$

Pronounced maxima in  $H_G(x,y,\theta,s)$  indicate that the image contains an object of interest centered at  $(x,y)$ , scaled by factor  $s$ , and rotated by angle  $\theta$ .


# Generalised Hough Transformation

- **Given:** the gradient vector  $(G_x, G_y)$  at each image location
- **Required:** possible locations of the object centre
  - How can possible centre locations be inferred from local gradient directions?
  - Which locations in  $H_G(x, y, \theta, s)$  are to be increased?




# Generalised Hough Transformation

- Exploit relation between local gradient and object center
- The local gradient is ambiguous
  - The same gradient occurs at different points on the object
  - A given gradient direction may correspond to several possible object centres
  - In general, several locations in  $H_G(x,y,\theta,s)$  must be increased



# Generalised Hough Transformation



- Exploit relation between local gradient and object center
- The local gradient is ambiguous
  - The same gradient occurs at different points on the object
  - A given gradient direction may correspond to several possible object centres
  - In general, several locations in  $H_G(x,y,\theta,s)$  must be increased






# Generalised Hough Transformation

- Exploit relation between local gradient and object center
- The local gradient is ambiguous
  - The same gradient occurs at different points on the object
  - A given gradient direction may correspond to several possible object centres
  - In general, several locations in  $H_G(x,y,\theta,s)$  must be increased




# Generalised Hough Transformation

- Exploit relation between local gradient and object center
- The local gradient is ambiguous
  - The same gradient occurs at different points on the object
  - A given gradient direction may correspond to several possible object centres
  - In general, several locations in  $H_G(x,y,\theta,s)$  must be increased




# Generalised Hough Transformation

- Exploit relation between local gradient and object center
- The local gradient is ambiguous
  - The same gradient occurs at different points on the object
  - A given gradient direction may correspond to several possible object centres
  - In general, several locations in  $H_G(x,y,\theta,s)$  must be increased



# Generating the R-Table

The R-Table stores, for each pixel on the edge of an object O, the local gradient direction and the displacement from the pixel to the object centre




## Row T of the R-Table

- Applies when local gradient direction is close to  $2\pi(T-1)/A$
- The displacement to the centre is  $(d_x^{T,1}, d_y^{T,1})$  or  $(d_x^{T,2}, d_y^{T,2})$  or ... or  $(d_x^{T,NT}, d_y^{T,NT})$

# Generating the R-Table

The R-Table stores, for each pixel on the edge of an object  $O$ , the local gradient direction and the displacement from the pixel to the object centre




$\theta$	
$[0, 30)$	$(2,1), (2,0), (2,-1), (-2,1), (-2,0), (-2,-1)$
$[30, 60)$	
$[60, 90)$	$(1,2), (0,2), (-1,2), (1,-2), (0,-2), (-1,-2)$
$[90, 120)$	
$[120, 150)$	
$[150, 180)$	

# Applying the R-Table

Considering an image location  $(x,y)$  with associated gradient direction  $\theta$ :

- $\theta$  is discretised by dividing the interval  $[0,2\pi]$  into  $A$  sub-intervals
- The sub-interval assigned to  $\theta$  identifies a row of the R-Table
- The row contains displacements  $(d_x^{T,1..NT}, d_y^{T,1..NT})$  to possible object centres
- $H_G$  can be increased by the local gradient magnitude at points

$$(x+d_x^{T,1}, y+d_y^{T,1}) \dots (x+d_x^{T,NT}, y+d_y^{T,NT})$$



# Applying the R-Table

Considering an image location  $(x,y)$  with associated gradient direction  $\theta$ :


- $\theta$  is discretised by dividing the interval  $[0,2\pi]$  into  $A$  sub-intervals
- The sub-interval assigned to  $\theta$  identifies a row of the R-Table
- The row contains displacements  $(d_x^{T,1..NT}, d_y^{T,1..NT})$  to possible object centres
- $H_G$  can be increased by the local gradient magnitude at points

$$(x+d_x^{T,1}, y+d_y^{T,1}) \dots (x+d_x^{T,NT}, y+d_y^{T,NT})$$


These steps compute  $H_G(x,y,\theta=0,s=1)$ .  
Scaled or rotated versions of the object are not detected.

# Generating the R-Table

The R-Table stores, for each pixel on the edge of an object  $O$ , the local gradient direction and the displacement from the pixel to the object centre




$\theta$	
$[0, 30)$	$(2,1), (2,0), (2,-1), (-2,1), (-2,0), (-2,-1)$
$[30, 60)$	
$[60, 90)$	
$[90, 120)$	$(1,2), (0,2), (-1,2), (1,-2), (0,-2), (-1,-2)$
$[120, 150)$	
$[150, 180)$	






# Generating the R-Table

The R-Table stores, for each pixel on the edge of an object  $O$ , the local gradient direction and the displacement from the pixel to the object centre




$\theta$	
$[0, 30)$	$(2,1), (2,0), (2,-1), (-2,1), (-2,0), (-2,-1)$
$[30, 60)$	
$[60, 90)$	
$[90, 120)$	$(1,2), (0,2), (-1,2), (1,-2), (0,-2), (-1,-2)$
$[120, 150)$	
$[150, 180)$	




# Generating the R-Table

The R-Table stores, for each pixel on the edge of an object  $O$ , the local gradient direction and the displacement from the pixel to the object centre





$\theta$	
$[0, 30)$	$(2,1), (2,0), (2,-1), (-2,1), (-2,0), (-2,-1)$
$[30, 60)$	
$[60, 90)$	
$[90, 120)$	$(1,2), (0,2), (-1,2), (1,-2), (0,-2), (-1,-2)$
$[120, 150)$	
$[150, 180)$	



# Generating the R-Table


The R-Table stores, for each pixel on the edge of an object  $O$ , the local gradient direction and the displacement from the pixel to the object centre



$\theta$	
$[0, 30)$	$(2,1), (2,0), (2,-1), (-2,1), (-2,0), (-2,-1)$
$[30, 60)$	
$[60, 90)$	
$[90, 120)$	$(1,2), (0,2), (-1,2), (1,-2), (0,-2), (-1,-2)$
$[120, 150)$	
$[150, 180)$	

# Changing the R-Table: Scale


- To detect objects of different scales, the R-Table is modified to describe a scaled version of the original object
- The updates to  $H_G$  are repeated at different scales to compute  $H_G(x,y,\theta=0,s)$




The displacements stored in the R-Table are multiplied by the desired scale factor before  $H_G$  is updated for the current scale

# Changing the R-Table: Scale

The R-Table stores, for each pixel on the edge of an object  $O$ , the local gradient direction and the displacement from the pixel to the object centre




$\theta$	
$[0, 30)$	$(2,1), (2,0), (2,-1), (-2,1), (-2,0), (-2,-1)$
$[30, 60)$	
$[60, 90)$	
$[90, 120)$	$(1,2), (0,2), (-1,2), (1,-2), (0,-2), (-1,-2)$
$[120, 150)$	
$[150, 180)$	




# Changing the R-Table: Scale

The R-Table stores, for each pixel on the edge of an object  $O$ , the local gradient direction and the displacement from the pixel to the object centre




$\theta$	
$[0, 30)$	$(2,1), (2,0), (2,-1), (-2,1), (-2,0), (-2,-1)$
$[30, 60)$	
$[60, 90)$	
$[90, 120)$	$(1,2), (0,2), (-1,2), (1,-2), (0,-2), (-1,-2)$
$[120, 150)$	
$[150, 180)$	




# Changing the R-Table: Scale

The R-Table stores, for each pixel on the edge of an object  $O$ , the local gradient direction and the displacement from the pixel to the object centre




$\theta$	
$[0, 30)$	$(2,1), (2,0), (2,-1), (-2,1), (-2,0), (-2,-1)$
$[30, 60)$	
$[60, 90)$	
$[90, 120)$	$(1,2), (0,2), (-1,2), (1,-2), (0,-2), (-1,-2)$
$[120, 150)$	
$[150, 180)$	




$\theta$	
$[0, 30)$	$(1,0.5), (1,0), (1,-0.5), (-1,0.5), (-1,0), (-1,-0.5)$
$[30, 60)$	
$[60, 90)$	
$[90, 120)$	$(0.5,1), (0,1), (-0.5,1), (0.5,-1), (0,-1), (-0.5,-1)$
$[120, 150)$	
$[150, 180)$	

# Changing the R-Table: Scale

The R-Table stores, for each pixel on the edge of an object  $O$ , the local gradient direction and the displacement from the pixel to the object centre



$\theta$	
$[0, 30)$	$(2,1), (2,0), (2,-1), (-2,1), (-2,0), (-2,-1)$
$[30, 60)$	
$[60, 90)$	
$[90, 120)$	$(1,2), (0,2), (-1,2), (1,-2), (0,-2), (-1,-2)$
$[120, 150)$	
$[150, 180)$	




$\theta$	
$[0, 30)$	$(1,0.5), (1,0), (1,-0.5), (-1,0.5), (-1,0), (-1,-0.5)$
$[30, 60)$	
$[60, 90)$	
$[90, 120)$	$(0.5,1), (0,1), (-0.5,1), (0.5,-1), (0,-1), (-0.5,-1)$
$[120, 150)$	
$[150, 180)$	



# Changing the R-Table: Orientation


- The R-Table is to be modified to describe a rotated version of the original object
- The rotation angle is assumed to be a multiple of  $\phi = (2\pi)/A$ 
  - Angular resolution in  $H_G$  is limited by the number of rows in the R-Table



- Gradients that used to be in row 1 are now associated with row 2
- The former last row becomes the new first row
  - Rotation causes a circular shift of the rows of the R-Table

# Changing the R-Table: Orientation

- The direction to the object center stays constant with respect to the local gradient




- The displacement vectors also rotate by angle  $\phi$ :


$$\begin{pmatrix} d_x^{a,b} \\ d_y^{a,b} \end{pmatrix} \leftarrow \begin{pmatrix} \cos(\phi) & -\sin(\phi) \\ \sin(\phi) & \cos(\phi) \end{pmatrix} \begin{pmatrix} d_x^{a,b} \\ d_y^{a,b} \end{pmatrix}$$

# Changing the R-Table: Orientation

The R-Table stores, for each pixel on the edge of an object  $O$ , the local gradient direction and the displacement from the pixel to the object centre




$\theta$	
$[0, 30)$	$(2,1), (2,0), (2,-1), (-2,1), (-2,0), (-2,-1)$
$[30, 60)$	
$[60, 90)$	
$[90, 120)$	$(1,2), (0,2), (-1,2), (1,-2), (0,-2), (-1,-2)$
$[120, 150)$	
$[150, 180)$	




# Changing the R-Table: Orientation

The R-Table stores, for each pixel on the edge of an object  $O$ , the local gradient direction and the displacement from the pixel to the object centre




$\theta$	
$[0, 30)$	$(2,1), (2,0), (2,-1), (-2,1), (-2,0), (-2,-1)$
$[30, 60)$	
$[60, 90)$	
$[90, 120)$	$(1,2), (0,2), (-1,2), (1,-2), (0,-2), (-1,-2)$
$[120, 150)$	
$[150, 180)$	

$$\begin{pmatrix} \cos(30) & -\sin(30) \\ \sin(30) & \cos(30) \end{pmatrix}$$




# Changing the R-Table: Orientation

The R-Table stores, for each pixel on the edge of an object  $O$ , the local gradient direction and the displacement from the pixel to the object centre




$\theta$	
$[0, 30)$	$(2,1), (2,0), (2,-1), (-2,1), (-2,0), (-2,-1)$
$[30, 60)$	
$[60, 90)$	
$[90, 120)$	$(1,2), (0,2), (-1,2), (1,-2), (0,-2), (-1,-2)$
$[120, 150)$	
$[150, 180)$	




$\theta$	
$[0, 30)$	
$[30, 60)$	$(1.2, 1.8), (1.7, 1), (2.2, 0.1), (-2.2, -0.1), (-1.7, -1), (-1.2, -1.8)$
$[60, 90)$	
$[90, 120)$	
$[120, 150)$	$(-0.1, 2.2), (-1, 1.7), (-1.8, 1.2), (1.8, -1.2), (1, -1.7), (0.1, -2.2)$
$[150, 180)$	

# Changing the R-Table: Orientation

The R-Table stores, for each pixel on the edge of an object  $O$ , the local gradient direction and the displacement from the pixel to the object centre




$\theta$	
$[0, 30)$	$(2,1), (2,0), (2,-1), (-2,1), (-2,0), (-2,-1)$
$[30, 60)$	
$[60, 90)$	
$[90, 120)$	$(1,2), (0,2), (-1,2), (1,-2), (0,-2), (-1,-2)$
$[120, 150)$	
$[150, 180)$	



$\theta$	
$[0, 30)$	
$[30, 60)$	$(1.2, 1.8), (1.7, 1), (2.2, 0.1), (-2.2, -0.1), (-1.7, -1), (-1.2, -1.8)$
$[60, 90)$	
$[90, 120)$	
$[120, 150)$	$(-0.1, 2.2), (-1, 1.7), (-1.8, 1.2), (1.8, -1.2), (1, -1.7), (0.1, -2.2)$
$[150, 180)$	


# Alternative Formulation using Correlation

- Traditional: **R-Table**



- The R-Table approach loops over pixels:
  - Consider each local gradient direction
  - Use displacement vectors from table to accumulate „votes“ in  $H_G$

- Alternative: **Direct Correlation**




- Explicitly search for the object in the image
- Loop over parameters  $(s, \theta)$ 
  - Scale  $s$  of the object
  - Orientation  $\theta$  of the object
  - Explicitly match the image with this version of the object
- The two formulations are (almost) identical


# Preprocessing the Example Object

The Object of interest  $O(x,y)$  is processed to obtain two representations

## 1. Extraction of the complex gradients $O^I(x,y)$



## 2. Obtain the binary edge image $O^B(x,y)$





- Binary image indicating pixels on edges  
→ Large local gradient magnitude
- Obtained using threshold  $T^B$ :

$$O^B(x, y) = |O^I(x, y)| > (T^B \cdot \max_{x, y} \{|O^I(x, y)|\})$$




# Preprocessing the Image

The image  $I(x,y)$  containing objects of interest is processed similarly to give complex gradients  $I^I(x,y)$



$$\frac{\partial}{\partial x} I(x,y)$$






$$\frac{\partial}{\partial y} I(x,y)$$

$$I^I = \frac{\partial}{\partial x} I(x,y) + i \cdot \frac{\partial}{\partial y} I(x,y)$$

# Correlation

$$H_G(x, y, \theta=0, s=1) = |\Re \{ (O^I(x, y) O^B(x, y)) \odot I^I(x, y) \}|$$




- The correlation is high where gradient directions in  $O$  and  $I$  match

$$K = \left| \sum_{o_j} |o_j| \cdot |i_j| \cos(\theta_j) \right|$$


- The cosine term means that votes can also become negative
- The absolute value ensures that detection is possible irrespective of whether the object is brighter or darker than the background

# Correlation in the Frequency Domain

- 1) Filter (object mask) and gradient image must have the same size before transformation  
→ Copy the object mask into a larger matrix



- 2) Centre the filter (the centre is shifted to coordinates (0,0) )



- 3) Transfer filter and image to the frequency domain
- 4) The spectrum of the result is the product of the image spectrum and the **complex conjugate** of the filter spectrum

# Scale and Orientation

The correlation is repeated with scaled and rotated object masks to fill  $H_G$

## 1) Changes in Scale

- Binary object mask and complex object gradients are resized

## 2) Rotation

- Binary object mask and complex object gradients are rotated
- This rotation does not affect the directions of complex gradients!
  - Complex gradients are rotated by applying a phase shift
  - After rotating by  $\theta$ , complex gradients are multiplied with  $\exp(-i \theta)$

## 3) Normalisation

- To retain scale invariance, the magnitudes of complex gradients are normalised:

$$O^I(x, y) \leftarrow \frac{O^I(x, y)}{\sum_x \sum_y |O^I(x, y)|}$$

# Exercise 3

- 1) Theory part
- 2) Implementing the transformation
- 3) Detecting multiple objects in an image

# Exercise 3 - Theory

The GHT provides the possibility to search for instances of the template object with different scale  $s$  and/or rotation  $\theta$  by using a 4D voting space  $H(x,y,s,\theta)$ .

The table on the right shows the R-table, which was generated by analysing the template object. As reference point the top-left corner of the object was chosen. The angles are defined relative to the x-axis and measured counter-clockwise.

Orientation	Displacement
[ 0, 45)	(0,10); (0,20)
[ 45, 90)	
[ 90,135)	
[135,180)	(-20,10); (-10,20)
[180,225)	
[225,270)	
[270,315)	(-10,0); (-20,0)
[315,360)	

## Task 1:

Which R-table below corresponds to scale ( $s=1$ ) with a counter-clockwise rotation of  $\theta = 90^\circ$ ?

i)		ii)		iii)		iv)	
Orientation	Displacement	Orientation	Displacement	Orientation	Displacement	Orientation	Displacement
[ 0, 45)	(-10,0); (-20,0)	[ 0, 45)	(0,-10); (0,-20)	[ 0, 45)		[ 0, 45)	
[ 45, 90)		[ 45, 90)		[ 45, 90)	(-10,0); (-20,0)	[ 45, 90)	(-20,10); (-10,20)
[ 90,135)		[ 90,135)	(-10,0); (-20,0)	[ 90,135)		[ 90,135)	
[135,180)	(-10,20); (-20,-10)	[135,180)		[135,180)		[135,180)	
[180,225)		[180,225)		[180,225)	(-10,-20); (-20,-10)	[180,225)	(-10,0); (-20,0)
[225,270)		[225,270)	(-10,-20); (-20,-10)	[225,270)		[225,270)	
[270,315)	(0,-10); (0,-20)	[270,315)		[270,315)		[270,315)	(0,10); (0,20)
[315,360)		[315,360)		[315,360)	(0,-10); (0,-20)	[315,360)	

# Exercise 3 - Theory

The GHT provides the possibility to search for instances of the template object with different scale  $s$  and/or rotation  $\theta$  by using a 4D voting space  $H(x,y,s,\theta)$ .

The table on the right shows the R-table, which was generated by analysing the template object. As reference point the top-left corner of the object was chosen. The angles are defined relative to the x-axis and measured counter-clockwise.

Orientation	Displacement
[ 0, 45)	(0,10); (0,20)
[ 45, 90)	
[ 90,135)	
[135,180)	(-20,10); (-10,20)
[180,225)	
[225,270)	
[270,315)	(-10,0); (-20,0)
[315,360)	


## Task 2:

Which simple geometric object does this R-table most likely represent?


- i) A triangle                      ii) A circle                      iii) A square                      iv) A parallelogram

# Given Functions

- **main** (Usage: `aia3 <path to template image> [<path to test image>]`)
  - Test routine for the generalised Hough transformation (main function)
  - Two different modes depending on number of input parameters:
    - 1<sup>st</sup> : Only template image is specified (one parameter)
      - Rotated and scaled version of template image is generated
      - Hough transformation is used to find template within this image
    - 2<sup>nd</sup> : Template as well as test image is specified (two parameters)
      - Hough transformation is used to find template within given image



**Object of  
Interest**




**Image**




# Given Functions


- **main** (Usage: `aia3 <path to template image> [<path to test image>]`)
  - Test routine for the generalised Hough transformation (main function)
  - Two different modes depending on number of input parameters:
    - 1<sup>st</sup> : Only template image is specified (one parameter)
      - Rotated and scaled version of template image is generated
      - Hough transformation is used to find template within this image
    - 2<sup>nd</sup> : Template as well as test image is specified (two parameters)
      - Hough transformation is used to find template within given image



**Object of Interest**




**Hough Space**




**Detection**

# Given Functions

- `main` (Usage: `aia3 <path to template image> [<path to test image>]`)
  - Test routine for the generalised Hough transformation (main function)
  - Two different modes depending on number of input parameters:
    - 1<sup>st</sup> : Only template image is specified (one parameter)
      - Rotated and scaled version of template image is generated
      - Hough transformation is used to find template within this image
    - 2<sup>nd</sup> : Template as well as test image is specified (two parameters)
      - Hough transformation is used to find template within given image



**Object of Interest**



**Hough Space**



**Image**

# Given Functions

- `main` (Usage: `aia3 <path to template image> [<path to test image>]`)
  - Test routine for the generalised Hough transformation (main function)
  - Two different modes depending on number of input parameters:
    - 1<sup>st</sup> : Only template image is specified (one parameter)
      - Rotated and scaled version of template image is generated
      - Hough transformation is used to find template within this image
    - 2<sup>nd</sup> : Template as well as test image is specified (two parameters)
      - Hough transformation is used to find template within given image

If an image contains more than a single object, a suitable thresholding procedure is needed to single out individual detections

## 1) Global Threshold

- Assume that all objects have received a vote in  $H_G$  that exceeds  $T^0 * \max(H_G)$
- $T^0$  is a fixed threshold

## 2) Identify Local Maxima


- Where an entire region of  $H_G$  received a high vote ( $>$ threshold), one should only consider local maxima of  $H_G$  as detection events.
- After applying the threshold, local maxima are considered the set of detected objects

# Given Functions


- `Mat makeTestImage(Mat& temp, double angle, double scale, double* scaleRange)`

- `temp`: Grayscale image
- `angle`: Rotation angle
- `scale`: Scaling factor
- `scaleRange`: Scale range

→ Generates a rotated and scaled version of the input image



**Object of  
Interest**



**Image**

# Given Functions

- `Mat rotateAndScale(Mat& temp, double angle, double scale)`


- `temp`: Grayscale image
- `angle`: Rotation angle in radian
- `scale`: Scale factor
- `return`: Transformed image

→ Scales and rotates given image

- `void circShift(Mat& in, Mat& out, int dx, int dy)`




- `in`: Image
- `out`: Shifted image
- `dx, dy`: Shift in x- and y-direction

→ Performs circular shift




# Given Functions

- `Mat calcDirectionalGrad(Mat& testImage, double sigma)`
  - `image`: Grayscale image
  - `sigma`: Standard deviation of filter (not necessarily integer)
  - `return`: Complex gradient image (two channel image)
- Computes the complex gradients in an image
- 1<sup>st</sup> channel: gradients in x-direction
- 2<sup>nd</sup> channel: gradients in y-direction



# Given Functions

- `void findHoughMaxima(vector< vector<Mat> >& houghSpace, double objThresh, vector<Scalar>& objList)`
    - `houghSpace`: The four dimensional Hough space  $H_G$
    - `objThresh`: Global threshold for object detection. Threshold is relative to the maximum of  $H_G$
    - `objList`: List of hough space coordinates (scale, angle, x, y)
- Applies a global threshold and identifies remaining local maxima in  $H_G$





# Given Functions

- `void plotHoughDetectionResult(Mat& testImage, vector<Mat>& templ, vector<Scalar>& objList, double scaleSteps, double* scaleRange, double angleSteps, double* angleRange)`
  - `testImage`: Grayscale image
  - `templ`: Object mask (see `makeObjectTemplate(..)`)
  - `objList`: Indices (into the Hough space  $H_G$ ) of the detected objects. Each entry specifies a set of coordinates: (scale, angle, x, y)
  - `scaleSteps`: Resolution of the Hough space  $H_G$  in the scale dimension
  - `scaleRange`: Minimum and maximum scales (e.g. `[0.5, 2.0]`)
  - `angleSteps`: Resolution of Hough space  $H_G$  in the orientation dimension
  - `angleRange`: Minimum and maximum angle (typically `[0, 2*pi]`)


→ Visualises the results of object detection.








# To Do

- `vector<Mat> makeObjectTemplate(Mat& templateImage, double sigma, double templateThresh)`
  - `templateImage`: Grayscale image with the object of interest
  - `sigma`: Std. deviation for the computation of image gradients
  - `templateThresh`: Threshold used to obtain the binary edge image (relative to the maximum gradient magnitude)
  - `return`: 2-element Mat-vector:
    - [0] contains the binary edge mask
    - [1] contains the corresponding complex gradients
- Creates a template of the object of interest, consisting of a binary edge map and a complex gradient image.



# To Do

- `void plotHough(vector< vector<Mat> >& houghSpace)`
  - `houghSpace`: Computed 4D hough space
- Creates a single image from the 4D hough space
- Displays and saves it as image file





# To Do

- `void makeFFTObjectMask(vector<Mat>& templ, double scale, double angle, Mat& fftMask)`
  - `templ`: Template for the object of interest (edge mask and gradients)
  - `scale`: New scale factor for the object
  - `angle`: New orientation for the object
  - `fftMask`: Spectrum of the mask that is correlated with a complex gradient image to find the object of interest, scaled by factor `scale` and rotated by `angle` radians.
- Computes the spectrum of  $O^I O^B$  for correlation with the complex gradient image.
  - Note:  $O^I O^B$  means **component wise multiplication**
- Scales and rotates the object template (updating the phase of gradients included), followed by normalisation, centering and Fourier transformation

# To Do

- `vector< vector<Mat> > generalHough(Mat& gradImage, vector<Mat>& templ, double scaleSteps, double* scaleRange, double angleSteps, double* angleRange)`
  - `gradImage`: Complex gradient image (with objects to be detected)
  - `templ`: Template of an object of interest
  - `scaleSteps`: Resolution of the Hough space in the scale dimension
  - `scaleRange`: Minimal and maximal scale factor (e.g. `[0.5, 2.0]`)
  - `angleSteps`: Resolution of the Hough space in the orientation dimension
  - `angleRange`: Smallest and largest orientation (typically `[0, 2*pi]`)
  - `return`: The computed Hough space  $H_G$ .  
Four “dimensions”: scale, orientation, spatial dimensions
- Computes the voting space of the generalised Hough transform
- **Be careful to note the sequence of dimensions!**
- Useful functions: `cv::mulSpectrum()`, `cv::dft()`

# Optional



- Apply GHT to other detection / recognition tasks  
→ eg. leaf detection & recognition task from last exercise

How does performance change? On what does it depend?